



Parallelization of Kempe's Influence Maximization Algorithm

EE451 Fall 2019

12/5/2019

Yi Sui, Pengcheng Li, Liheng Lin



Introduction to Influence Maximization

- Social Network & “Influencer”
- Spread of Information, ideas and influence
 - The use of new product
 - A new government policy
 - News
- Maximization of Information, ideas and influence
 - Different influence from one to another
 - Choose key individuals(best influencers)




Influence Maximization

- Weighted directed graph $G = (V, E, w)$
 - $0 \leq w \leq 1$
 - Weight: probability that a vertex can influence another
- Influence of a vertex v
 - A vertex influences neighbors with given probability
 - Might succeed to influence or fail
 - Successfully Influenced vertices continue to influence their neighbors
 - $\sigma(v)$ = The number of vertices being influenced



Graph





Influence Maximization

- Influence of a set S
 - Every vertex in S have an influenced set
 - $\sigma(S)$ = The size of union of influenced sets
- Influence Maximization:
 - Given a maximum set size k
 - Return a set of size k with maximum influence $\sigma()$



Influence Maximization Problem

- NP-hard
 - Even an approximate algorithm runs for days when there are billions of vertices!
- Large graph
 - May not fit in the memory

Need parallelization to solve the time and space issue!



Kempe's Influence Maximization Algorithm

- Approximation Algorithm
- Greedy Strategy
 - Maintain a seed set
 - In each step, add a vertex that helps maximize the influence
 - Stop until maximum seed size is achieved



Kempe's Influence Maximization Algorithm

```
S =  $\emptyset$ 
while |S| < k:
    for each vertex v in V-S:
        simulate  $\sigma(S \cup \{v\})$  by sampling for a large number of
        times
    end for
    v* = argmax  $\sigma(S \cup \{v\})$ 
    S = S  $\cup$  {v*}
end while
return S
```

This algorithm needs a large amount of computation for a mid-size graph



Kempe's Influence Maximization Algorithm

$S = \emptyset$

while $|S| < k$:

→ for each vertex v in $V-S$:

→ simulate $\sigma(S \cup \{v\})$ by sampling for a large number of

times

Independent Work: Parallelizable!

$S = S \cup \{x\}$

end while

return S



Parallelization

3 parallelization approaches:

- Static Mapping
- Dynamic Mapping
- Graph Partitioning



Mapping

$S = \emptyset$

while $|S| < k$:

→ for each vertex v in $V-S$:

 simulate $\sigma(S \cup \{v\})$ by sampling for a large number of
 times

 end for

$v^* = \operatorname{argmax} \sigma(S \cup \{v\})$

$S = S \cup \{v^*\}$

end while

return S



MPI - Static Mapping

- K vertices and P processes, each process is responsible for a block of K/P vertices
- Each process stores the entire graph locally
- Parallel:
 - Each process begins simulation from one vertex v in its block
 - Calculate average influence of v
 - Find out v with max influence in its block
- Gather P vertices with max avg influence from P processes
- Add the vertex with global max influence into Seed



MPI - Dynamic Mapping

- Hypothesis:
 - Imbalanced workload due to different probability
 - Dynamic Mapping for balanced workload to reduce runtime
- Master-worker model
- Instead of evenly dividing the vertices, master assign new vertices(a block) to workers whenever they finish current work
- Communication cost increases but idle time decreases



MPI - Graph Partitioning

- K vertex and P processes, each process is responsible for a block of K/P vertices
- Only read the assigned part of graph into memory
- Communicate when influencing a vertex stored in other processes
- Communication time increases but memory access time decreases
- Better performance in larger graphs



Platform

- MPI Model
 - Suitable for communication
 - **Mapping**: distribute influence simulations to workers
 - Graph-partition: not memory sharing model
 - While loop: not for OpenMP



Challenges

- Unfamiliar with MPI non-blocking send/receive
 - Synchronization issues
- Graph Partitioning
 - Partition & Synchronization
- HPC(testing & debugging)
 - Slow resources allocation



Raw Dataset

- *Bitcoin OTC trust* network example from the Stanford Large Network Dataset Collection
 - A trust network of Bitcoin investors from the online trade platform Bitcoin OTC
 - Buyer & seller can rate each other after a transaction
 - range: ± 10 (edge weight)
- Weighted signed directed temporal network
- 5,881 vertices and 35,592 directed edges
- 89% edges are positive



Raw Dataset

```
source,target,rate,timestamp
6,2,4,1289241911.72836
6,5,2,1289241941.53378
1,15,1,1289243140.39049
4,3,7,1289245277.36975
13,16,8,1289254254.44746
13,10,8,1289254300.79514
7,5,1,1289362700.47913
2,21,5,1289370556.80938
2,20,5,1289370622.21473
21,2,5,1289380981.52787
21,1,8,1289441411.46365
21,10,8,1289441438.4426
21,8,9,1289441450.59006
```



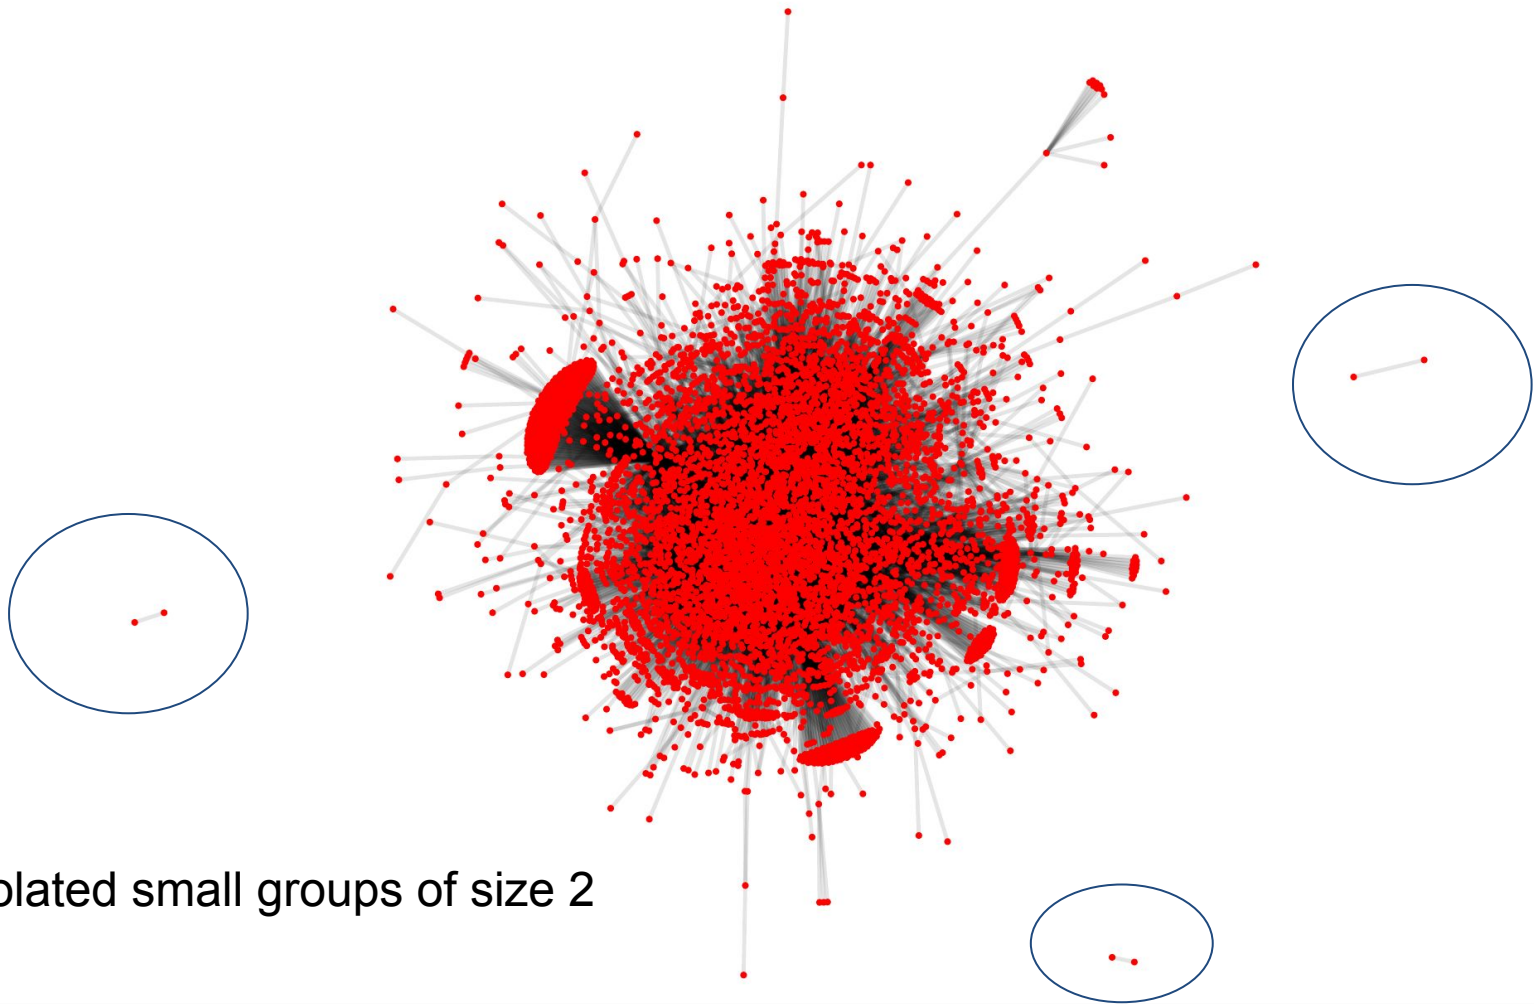
Dataset

- Processed into a weighted directed network
- 5,881 vertices and 35,592 directed edges
- Shift & normalize edge weight(rate) from ± 10 to $[0, 1]$

```
source,target,weight  
1,2,0.9  
1,3,0.8  
1,4,1.0  
1,5,0.7  
1,6,0.9
```



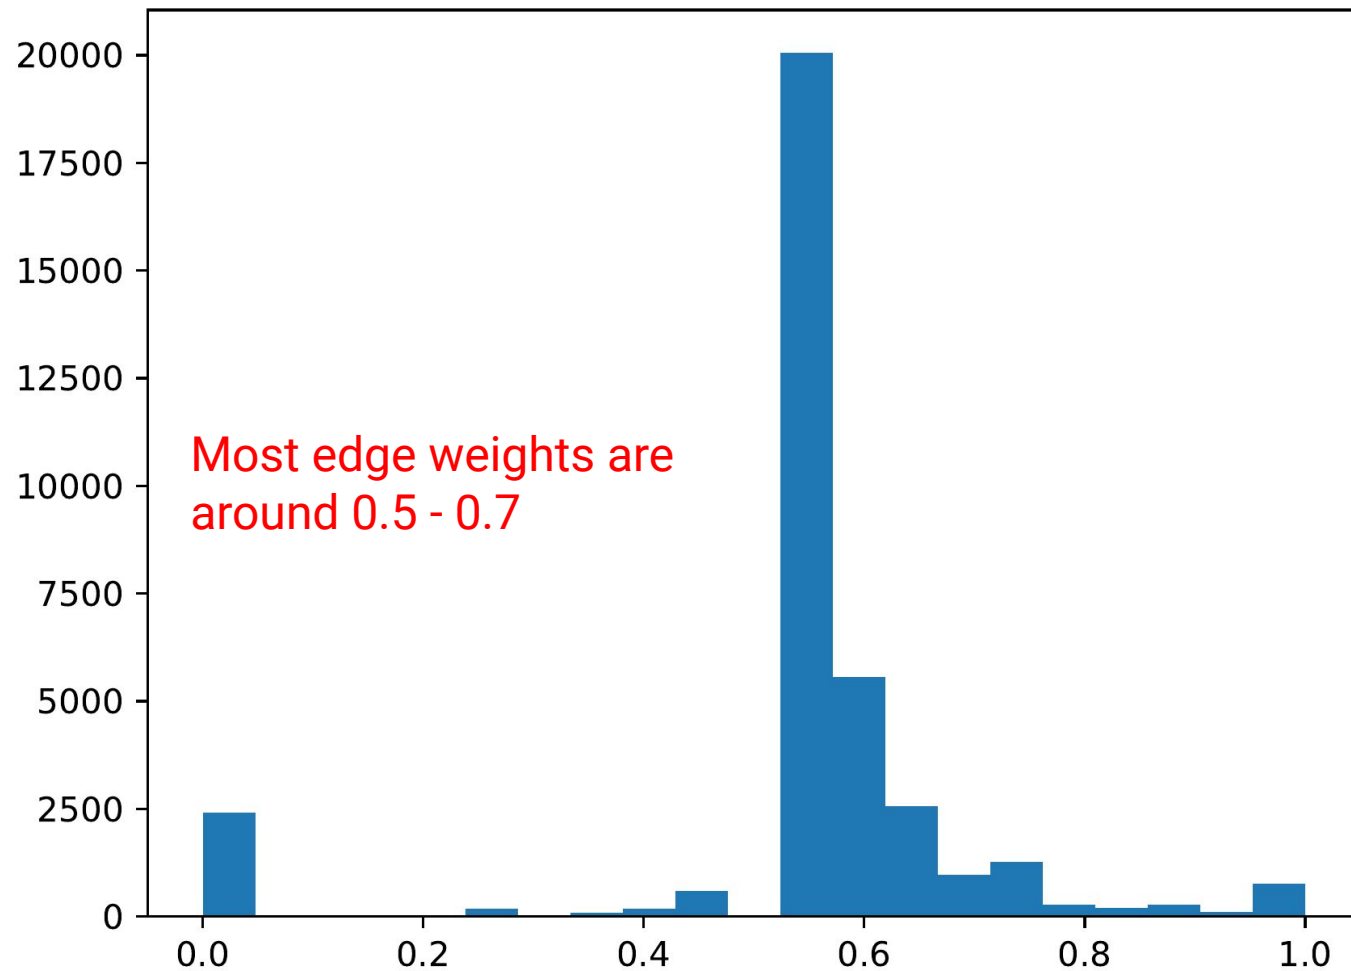
Dataset - Graph



3 isolated small groups of size 2



Dataset - Distribution





Explanation of Parameter

- Seed size(Input parameter)
 - How many vertices to pick and return?
- Sample Times(Input parameter)
 - Repeat the influence process to improve accuracy
- Block size in Dynamic Mapping(parallelization parameter)
 - Balancing workload vs communication cost
- Number of workers(parallelization parameter)
 - Using 2 vs 4 workers



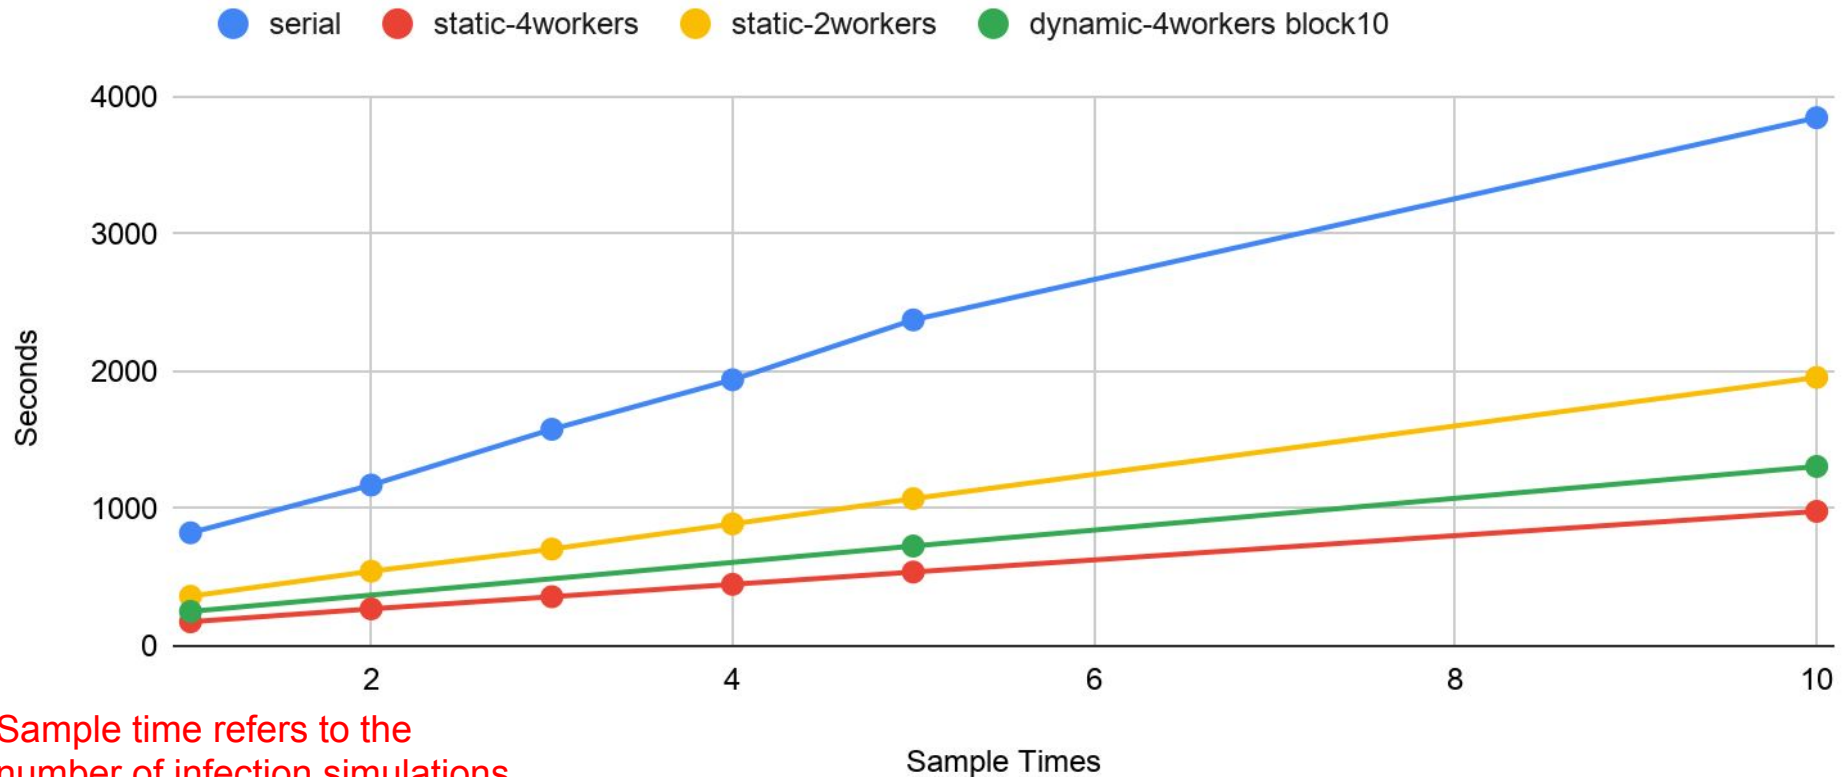
Testing Platform

- USC HPC
 - Serial: allocate 1 computing node
 - Parallel: allocate 2/4 computing nodes
 - `#SBATCH --ntasks-per-node=1`
 - `#SBATCH --nodes=1/2/4`



Results

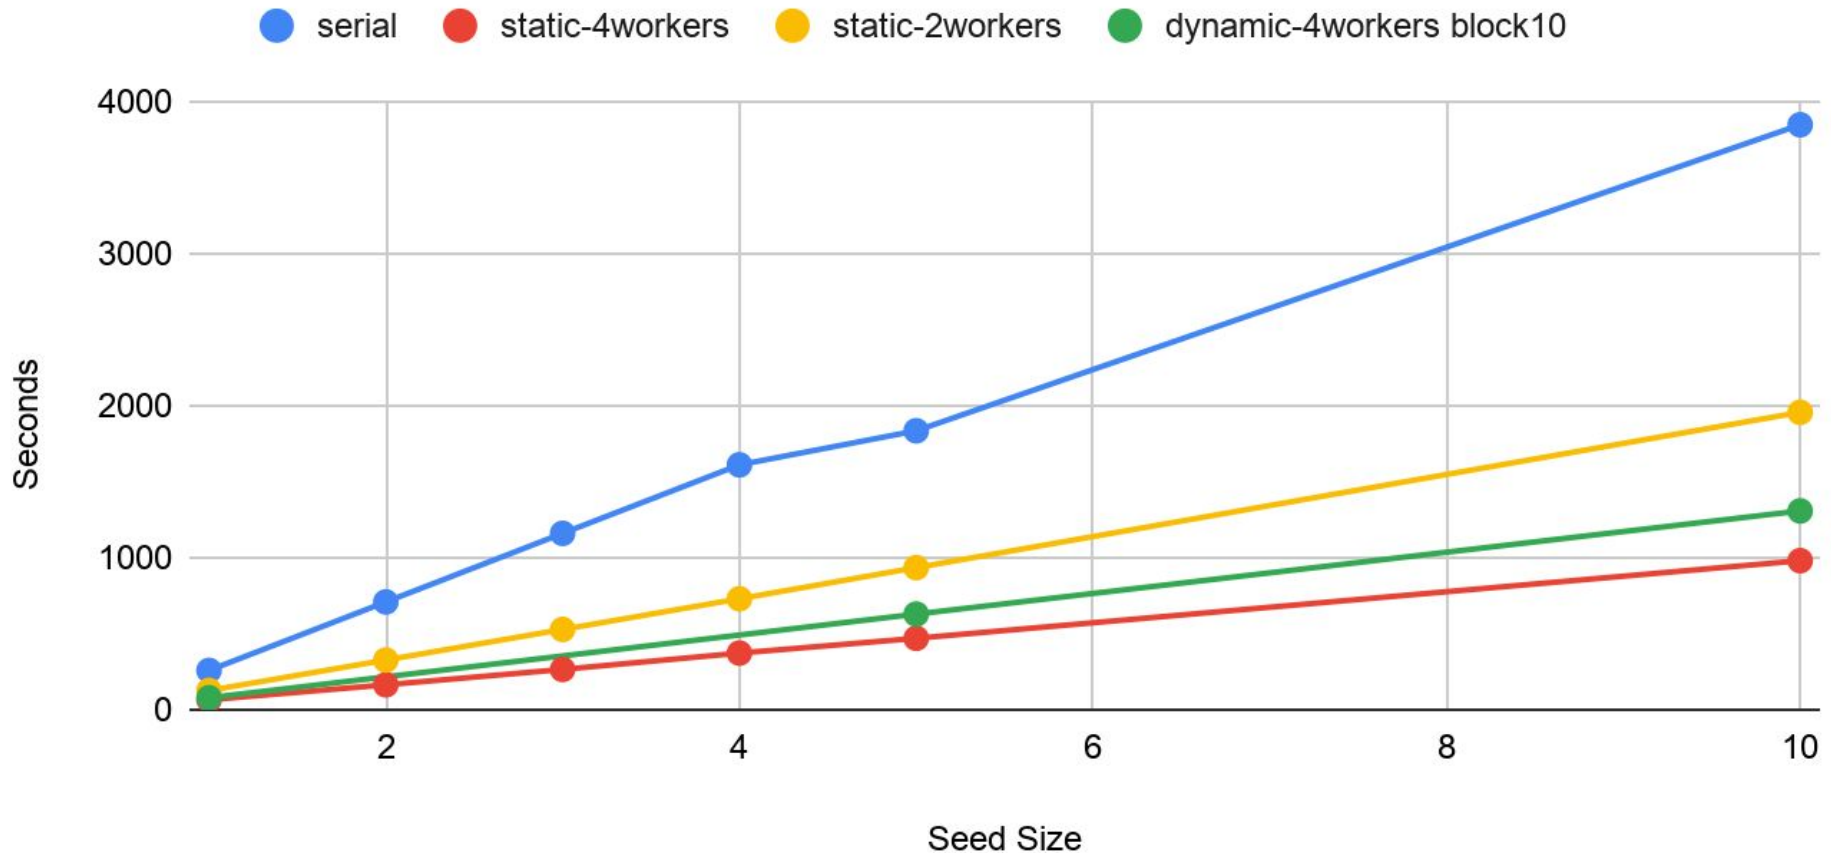
Average Time(3 runs) vs Different Sample Times, Seed Size =10





Results

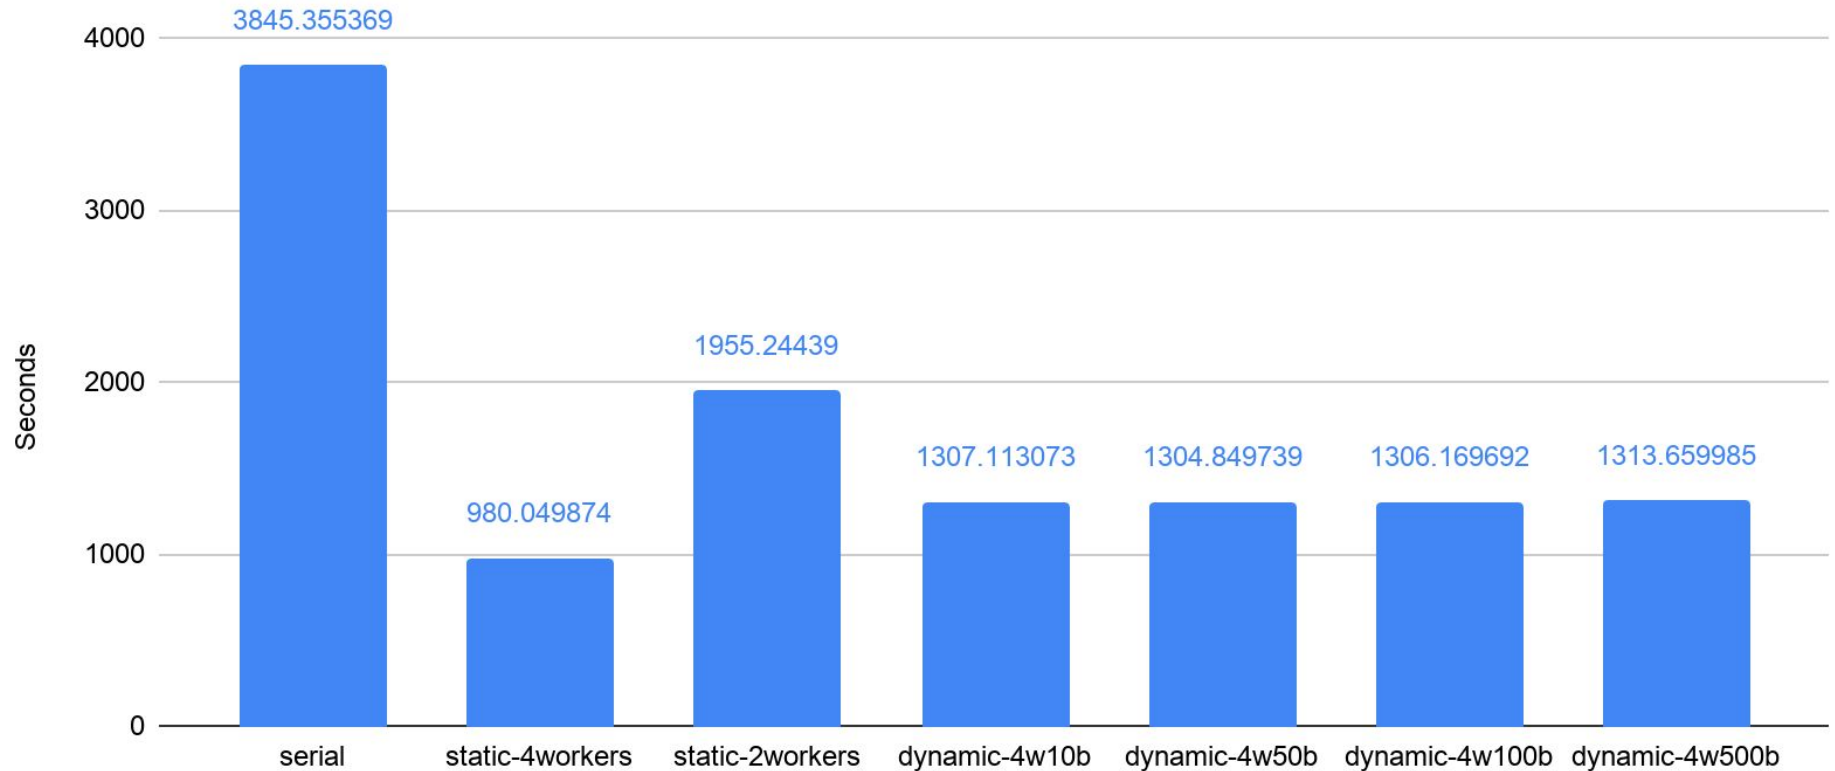
Average Time(3 runs) vs Different Seed Sizes, Sample Time =10





Results

Average Time(3 runs) vs Different Implementations
Seed Size = 10, Sample Times = 10





Result - Graph Partitioning

- Implementation is complete
 - Results are correct, but performance can be improved
 - Communication & synchronization costs
 - Barrier for synchronization
 - Influence steps moving between vertices in different processes - extra communication



Analysis

- Both Static Mapping and Dynamic Mapping are faster than the serial version
 - Static w/ 4 workers: 3.9x Speedup
 - Static w/ 2 workers: 2.0x Speedup
- Dynamic w/ 4 workers: 2.9x Speedup



Analysis

- Dynamic Mapping is slower than Static Mapping
 - One process serves as Master
 - Master does not compute
- Dynamic Mapping with different block sizes:
 - Time difference was ignorable
 - Communication cost doesn't affect much in testing
 - Possible causes:
 - Workload is not very unbalanced
 - Graph is not large enough to show the difference



Conclusion

- We have implemented:
 - Serial version of an Influence Maximization Problem
 - Several parallelized versions
- We learn...
 - MPI Programming



Question?